1 Introduction

We present the translation from a simple high-level language with functional elements and LISP-like syntax to SPIR-V, an abstract assembly language for GPU shaders and kernels.

The idea of using functional concepts in a shader came from the fact that shaders usually have a small and well-defined set of inputs and outputs and can therefore mostly be written as pure computations.

Originally, another personal motivation for developing a shader language translatable to SPIR-V was that existing shader languages have some properties that are undesirable in certain contexts. For instance, to optimize shaders for performance many operations (read: programming errors) like accessing an array or input out-of-bounds, but even normalizing a zero-length vector or any computations involving infinity are undefined in certain contexts and environments. In practice, this undefinedness is often unobservable in some implementations (leading to no error) while it leads to major errors on other implementations, making development and debugging for all implementations hard. The idea was to develop a language that guarantees to make undefined behavior observable (e.g. by always checking for undefined behavior conditions at runtime and write output to a buffer should they be fulfilled). This obviously comes at a potentially huge and unacceptable runtime cost and brings no advantage for shipped application but could be quite useful in a development environment, making sure that shader programs don't trigger undefined behavior. During the development of this language I realized though that this is better implemented as a completely separate SPIR-V-only pass since it doesn't depend at all on anything in the source language. Regarding the correctness proof of the compiler it is furthermore not too interesting, since it would just introduce additional branches for builtin operations. The following sections hint at it a couple of times but in the end I only describe the translation from our source language to SPIR-V, adopting their (or the runtime's) notions of undefined behavior.

During this project I found that SPIR-V does not have defined operational semantics and even the statements regarding undefined behavior (which are mainly specified by the runtime environments like Vulkan or OpenCL and not SPIR-V itsefl) weren't completely unambiguous.

1.1 Example

The most simple valid program we can write is a fragment shader (i.e. the type of program executed for each pixel [1] of a rasterized primitive) just writing a constant color. In our shader language it just looks like this, writing the color opaque yellow in this case:

```
( vec4  1.0  1.0  0.0  1.0)
```

It translates to the following program by our compiler:

```
        OpCapability  Shader
%3 =  OpExtInstImport  "GLSL.std.450"
        OpMemoryModel  Logical  GLSL450
        OpEntryPoint  Fragment  %1  "main"  %16
        OpExecutionMode  %1  OriginUpperLeft
        OpDecorate  %16  Location  0
```

---

[1]This is an oversimplification but enough for the basic understanding

```
    %float = OpTypeFloat 32
     %void = OpTypeVoid
   %v4float = OpTypeVector %float 4
     %bool = OpTypeBool
        %2 = OpTypeFunction %void
     %true = OpConstantTrue %bool
    %false = OpConstantFalse %bool
   %float_1 = OpConstant %float 1
 %float_1_0 = OpConstant %float 1
   %float_0 = OpConstant %float 0
 %float_1_1 = OpConstant %float 1
%_ptr_Output_v4float = OpTypePointer Output %v4float
       %16 = OpVariable %_ptr_Output_v4float Output
        %1 = OpFunction %void None %2
       %10 = OpLabel
       %15 = OpCompositeConstruct %v4float %float_1 %float_1_0 %float_0 %float_1_1
            OpStore %16 %15
            OpReturn
            OpFunctionEnd
```

A huge part of this is header and wrapper code we always generate (and constants, not part of the main program). In this case, the only interesting instruction (besides the defined constants) is the creation of the vector value via OpCompositeConstruct into SPIR-V ID 15, everything else is (more or less) always generated like this.

2   Source language: λV

Our source language is a simple shader language with functional elements. To keep proofs and formalization minimal, we will only look at a simple version of λV (basically already a subset of what I implemented [2]):
We only look at fragment shaders and only allow a single observation: one output to the framebuffer. This isn't too much of a simplification for realistic fragment shaders and modeling writes to multiple framebuffer attachments as well as inputs shouldn't be too difficult but require some complications to our formalization. Similarly, there isn't a huge difference to other shader types (such as vertex or compute shaders) except that those usually produce more or different observations and therefore one just once again needs to model the observations in a more complicated way (especially when allowing arbitrary buffer or image stores as needed to make compute shaders useful). But I don't expect the simplified model we use here to be too hard to extend to cover the given cases.

The syntax of λV is as simple as possible. Expressions can be numbers, true or false, an identifier or a list (consisting of zero or more expressions). There would be no advantage in encoding builtins such as $+$, $func$ or if into the syntax of the language. Those are simply pre-defined identifiers. In practice, keeping the syntax this simple has the advantage that writing a parser and AST representation is extremely trivial.

---

[2] See https://github.com/nyorain/lambdaV for my experiments with implementing a real compiler, a lot of the formal proof here is actually quite close to the compiler source code. It supports multiple outputs, and define statements, i.e. a program can consist of more than one expression

$$e ::= num \,|\, true \,|\, false \,|\, identifier \,|\, (l)$$
$$l ::= \varnothing \,|\, e\,l$$

3   The target language: SPIR-V

SPIR-V [3] is an SSA-form [4] abstract typed assembly language. It has similarities to LLVM IR but is more limited: no dynamic dispatch, i.e. no function pointers, no memory allocations, no stack frame and most runtimes (i.e. where SPIR-V modules can be used as shaders or kernels) don't allow recursion, see for instance the Vulkan specification[5] for this. On the other hand, SPIR-V provides some GPU-specific primitives. It's specification does not give operational semantics - or any formal specification - at all but rather describes the layout and semantics in plain text. To formally argue about correctness we need to model at least some form of operational semantics though. We will only consider the subset of SPIR-V used by our compiler. For instance, our compiler never outputs any functions (apart from the main entrypoint) so we don't care for that.

We define SPIR-V execution state like this $S := (M, \vec{I}, ID, ID)$
ID is a SPIR-V identifier, i.e. just a number. M is of the memory of form $[id \mapsto V]...$ and maps IDs to values. Values can be

- instruction blocks of form $\vec{I}$

- values of form (V, τ), where V the first value is a a primitive (number, vector, matrix, bool) and τ a type for this value.

$\vec{I}$ is an instruction vector. Furthermore, our state contains the current and previous block IDs, this is needed to resolve SSA phi [6] instructions.

Our judgment looks like this: $S \to S$.

We formally model execution of a (valid) SPIR-V module like this: when the module is loaded, all constants and types (declared in the header) are loaded into a memory M. All instruction blocks are loaded into the memory as well, removing the first $OpLabel$ instruction that is only used to identify the blocks with an id. The header of the SPIR-V module defines the entry point function. Execution looks up the first block in the entry point function (functions must start with a label defining the first block id) with instruction vector $\vec{I}_{entry}$ and block ID id_{entry} and then behaves as specified by the operational semantics for $M, \vec{I}_{entry}, id_{entry}, 0$.

We represent SPIR-V in the standard textual assembly format, with the new ID defined by the instruction (if any) on the left of the "=" sign.

---

[3]https://www.khronos.org/registry/spir-v/specs/unified1/SPIRV.html

[4]See https://en.wikipedia.org/wiki/Static_single_assignment_form

[5]https://www.khronos.org/registry/vulkan/specs/1.2/html/chap35.html#spirvenv

[6]See again https://en.wikipedia.org/wiki/Static_single_assignment_form

3.1   Control flow

$$\frac{M(tid) = \vec{I_t}}{M, (\text{OpBranch } id_t), c, p \to M, \vec{I_t}, id_t, c}$$

$$\frac{M(id_c) = (true, Bool) \qquad M(id_t) = \vec{I_t}}{M, (\text{OpBranchConditional } id_c \ id_t \ id_f), c, p \to M, \vec{I_t}, id_t, c}$$

$$\frac{M(id_c) = (false, Bool) \qquad M(id_f) = \vec{I_f}}{M, (\text{OpBranchConditional } id_c \ id_t \ id_f), c, p \to M, \vec{I_f}, id_f, c}$$

$$\frac{p = parent_i \qquad M(var_i) = V \qquad M' = M[id_r \mapsto V]}{M, (id_r = \text{OpPhi } id_{type} \ var_1 \ parent_1 \ldots var_n \ parent_n, \vec{I}), c, p \to M', \vec{I}, c, p}$$

3.2   Computations

Most instructions simply run a computation. Defining the operational semantics for those is not too interesting and defining all the rules for the various functions wouldn't be helpful. They all look more or less like this example we give for floating point addition:

$$\frac{M(id_1) = (num_1, Num) \qquad M(id_2) = (num_2, Num) \qquad M' = M[id_r \mapsto fadd(num_1, num_2)]}{M, (id_r = \text{OpFAdd } id_{type} \ id_1 \ id_2, \vec{I}), c, p \to M', \vec{I}, c, p}$$

where $fadd$ simply encodes the semantics of the addition of two numbers. Interestingly enough, SPIR-V does not specify how overflow or special cases (infinity or NaN arguments) are handled. Instead, this is usually specified in the runtime environment. We can for instance look once again into the Vulkan specification. It specifies that "By default, the implementation may perform optimizations on half, single, or double-precision floating-point instructions that ignore sign of a zero, or assume that arguments and results are not NaNs or infinities." and "NaNs may not be generated. Instructions that operate on a NaN may not result in a NaN.". Basically everything that uses or results in special floating point values is undefined behavior, more or less. Newer SPIR-V versions support a flag signaling that infinities and NaNs must be preserved (since SPIR-V 1.4) but we also want to target SPIR-V 1.0 and implementations that do not provide the optional support for this flag. Given these non-guarantess, we can't even check for infinity or NaN *after* we do an operation since then we might already have triggered undefined behavior, at least that is my interpreation of this specification. But checking whether an operation might overflow (or similar) is a pain (maybe not even possible, given that Vulkan and SPIR-V give their implementations some freedom regarding rounding of values returned by computations). Apparently, I'm also not the first person confused by this, see Vulkan-Docs issue 961.

Possible solutions to the problem include these:

- Just make any overflow or similar undefined in the source language as well. I don't want to do this since one of the main motivations in the first place was to get an target language program that is as deterministic as possible (at least detecting triggered undefined behavior).

- Actually evaluating whether an operation would operate on or return infinity or NaN at runtime. For each operation. That's such a huge pain. I'm not even thinking about runtime cost here, implementing a check that safely evaluates whether addition (and multiplication, division, exp, exp2, pow, ...) of floating point numbers would give such a result seems like a lot of work and definitely not a sane solution to me.

- Just outputting SPIR-V 1.4 (or using the previously available extension) and requiring support for this flag. That is what is assumed in the rest of this report and what will probably be used in the further implementation of the compiler since my hardware supports it and it makes a big issue just go away, basically. Most computations don't ever trigger any undefined behavior anymore.

There are still some instructions, however, that trigger undefined behavior. As mentioned in the introduction, we would optimally like to have *debug* mode of binaries in which triggering undefined behavior is somehow signaled, helping the programmer to remove it from programs, even if the implementation they are using just happens to implement the undefined behavior in a way that does not result in any problems. But for now, our source language just inherits the undefined behavior conditions of SPIR-V. In our conclusion we go into some detail on a better way to make undefined behavior observable in SPIR-V in general.

3.3   Ignored instructions

There are meta-instructions that must be inserted into a SPIR-V module for correcntess that don't have an impact on the semantics, like the OpSelectionMerge and OpLoopMerge instructions that provide meta-information about the control flow. Furthermore there are debug instructions, allowing to associate source-language line numbers or variables names with SPIR-V code. We simply ignore all those instructions, treating them the same way we treat OpNop:

$$\overline{M, (\text{OpNop}, \vec{I}), c, p \rightarrow M, \vec{I}, c, p}$$

4 λV operational semantics

To reason about correctness properties of our translation, we will define small-step operational semantics for λV.

$$\overline{(let\ ((id_1\ e_1) \ldots (id_n\ e_n))\ e) \rightarrow e[e_1/id_1] \ldots [e_n/id_n]}$$

$$\overline{(if\ true\ e_2\ e_3) \rightarrow e_2}$$

$$\frac{}{(if\ false\ e_2\ e_3) \to e_3}$$

$$\frac{}{((func\ (id_1 \ldots id_n)\ e)\ e_1 \ldots e_n) \to e[e_1/id_1] \ldots [e_n/id_n]}$$

$$\frac{}{((\text{rec-func}\ (id_1 \ldots id_n)\ e)\ e_1 \ldots e_n) \to e[e_1/id_1] \ldots [e_n/id_n][(\text{rec-func}\ (id_1 \ldots id_n)\ e)/rec]}$$

The reduction rule for builtins are intuitively defined, just copying the underlying SPIR-V (and target environment, e.g. Vulkan) semantics, which in turn usually just refer to the IEEE floating point standard.

The design decision we made for handling the (by default ill-defined) issues like overflow, infinities and NaN's in Section 3.2 is important here since it significantly modifies the semantics of those builtins in SPIR-V. Otherwise we would transitively introduce a lot of undefined behavior in our source language.

The conversion semantics are fairly simple due to our general list syntax:

$$\frac{\vdash e \to e'}{\vdash e \to^+ e'}$$

$$\frac{\vdash e_0 \to^+ e_0'}{\vdash (e_0 \ldots) \to^+ (e_0' \ldots)}$$

Reduction semantics:

$$\frac{}{\vdash e \to^* e}$$

$$\frac{\vdash e \to^+ e'' \qquad \vdash e'' \to^* e'}{\vdash e \to^* e'}$$

As usually, we define a evaluation function, $eval_\lambda(e) = o$, that is defined as the observation o that e can be reduced to. All possible observations in our simplified version of the source language is a single $Vec(I, f)$, the value written to the framebuffer by the fragment shader.

5   Types in λV

$\tau ::= Num \mid Bool \mid Vec(\{2, 3, 4\}, \{f, b\}) \mid Mat(\{2, 3, 4\}, \{f, b\}) \mid Rec\ \tau \mid PureRec$

Num is a type for all numbers. We don't seperate between integers and floating point numbers, we just assume all numbers to be floating-point for simplicity and since GPUs are usually optimized for that anyways. For logical true and false, we have Bool. Then, we have various Vector and Matrix types (with dimensions 2, 3 or 4 and floating point "f" or boolean "b" elements). For simplicity, we only allow square-sized matrices at the moment, other ones are rarely needed in shaders anyways. The types $Rec\ \tau$ and $PureRec$ have nothing to do with recursive types but are just helper types that allow us to model the restriction we have to put on recursive functions (namely: only tail-recursion is allowed) while still deducing the types of recursive expressions. This is probably the only interesting thing about the type

system. How exactly the types are used should become apparent from the typing rules below. To allow minimizing the number of rules and distinct cases, we will write $Vec(1, \{b, f\})$ as synonym for *Bool* or *Num*, respectively.

 Our typing judgment has the following form:

$$\boxed{R, A \vdash e : \tau}$$

R is a *recursive context*, holding the types of the arguments in the current recursive frames, $R ::= \tau*$. A is a tuple of tuple of expressions, representing the current stack of call arguments, $A ::= (e*)*$. Both R and A are basically needed as workaround for not typing functions while allowing them in almost any context (there are some technical limitations discussed below). A typing judgement means that e is of type τ, with arguments A and in recursive context R. We also use T, U as type metavariables. [I guess this is fairly common but for tuples we write $((a_1 \ldots), a_2)$ for $(a_1 \ldots a_2)$, i.e. appending to a tuple].

$$\text{t-num} \quad \frac{}{R, \varnothing \vdash num : Num}$$

$$\text{t-true} \quad \frac{}{R, \varnothing \vdash true : Bool}$$

$$\text{t-false} \quad \frac{}{R, \varnothing \vdash false : Bool}$$

$$\text{t-if} \quad \frac{\varnothing, \varnothing \vdash e_1 : Bool \qquad R, A \vdash e_2 : tau_1 \qquad R, A \vdash e_3 : tau_2}{R, A \vdash (if\ e_1\ e_2\ e_3) : \text{rec-match}(tau_1, tau_2)}$$

$$\text{t-app} \quad \frac{R, (A, (e_1 \ldots e_n)) \vdash e_0 : \tau}{R, A \vdash (e_0\ e_1 \ldots e_n) : \tau}$$

$$\text{t-func} \quad \frac{R, A_r \vdash e[e_1/id_1] \ldots [e_n/id_n] : \tau}{R, (A_r, (e_1 \ldots e_n)) \vdash (func\ (id_1 \ldots id_n)\ e) : \tau}$$

$$\text{t-rec-func} \quad \frac{R, \varnothing \vdash e_i : \tau_i \qquad \vec{\tau_i}, \varnothing \vdash e[e_1/id_1] \ldots [e_n/id_n] : Rec\ \tau}{R, ((e_1 \ldots e_n)) \vdash (\text{rec-func}\ (id_1 \ldots id_n)\ e) : \tau}$$

$$\text{t-rec} \quad \frac{n > 0 \qquad len(\tau_i) = n \qquad \varnothing, \varnothing \vdash e_i : \tau_i}{\vec{\tau_i}, ((e_1 \ldots e_n)) \vdash rec : PureRec}$$

$$\text{t-let} \quad \frac{R, A \vdash e[id_1/e_1] \ldots [id_n/e_n] : \tau}{R, A \vdash (let\ ((id_1\ e_1) \ldots (id_n\ e_n))\ e) : \tau}$$

Substituion is assumed to be context-sensitive (i.e. only substitute those identifiers that are really meant in that case and not those that are redefined in a deeper scope), as usually.

The helper function *rec-match* combines two related types in a recursive context.

$$rec-match(\tau_1, \tau_2) := \begin{cases} \tau_1, & \text{for } \tau_1 = \tau_2, \\ \tau_1, & \text{for } \tau_1 = Rec\,\tau_2, \\ \tau_2, & \text{for } \tau_2 = Rec\,\tau_1, \\ \tau_2, & \text{for } \tau_1 = PureRec \land \tau_2 = Rec\,T, \\ \tau_1, & \text{for } \tau_2 = PureRec \land \tau_1 = Rec\,T, \\ Rec\,T, & \text{for } \tau_2 = PureRec \land \tau_1 = T \text{ (where T isn't } PureRec \text{ or } Rec\,U), \\ Rec\,T, & \text{for } \tau_1 = PureRec \land \tau_2 = T \text{ (where T isn't } PureRec \text{ or } Rec\,U) \end{cases}$$

PureRec is the type of a *rec* call but when we have an *if* expression where one branch just results in a *rec* call, i.e. having type *PureRec* while the the other branch contains a value of type τ (or $Rec\,\tau$), we can deduce that the function must return a value of type τ in general. But instead of giving this *if* expression then the type τ, we give it the type $Rec\,\tau$ since the returned value can't be used for any further computations (except control flow, at the moment this only means *if*). Furthermore, this type system encodes the requirement for *rec-func* constructs to have at least one *rec* call in its body (since typing requires the body of *rec-func* to be of type $Rec\,\tau$).

When typing expressions inside a *rec-func* construct, the recursive context R holds the types the *rec-func* was called with. This means that one cannot call recursive functions with function objects (since they are not typed). In practice this is a technical limitation we cannot overcome since SPIR-V does not support dynamic dispatch and allowing recursion on arbitrary (possibly different for each recursive call) functions yields cases where we can't inline function calls anymore, i.e. can't unroll recursive functions to simple loops. In practice, we could put a more relaxed restriction on our type system: It is allowed to call recursive functions with function values as long as all recursive calls use the same function value. Or even more general: As long as there is only a finite number of functions used in the recursive calls (meaning basically that you recurse with newly instantianted functions in each recursion, it should be somewhat intuitive that a case like that can't be inlined/unrolled anymore). But that is a much more complicated restriction, yields a more complex type system and code generation. And in practice one can simply use workarounds. For instance:

```
(let ((nat-fold (func (n accum f) (
  let ((body (rec-func (n accum) (
    if (eq n 0)
      accum
      (rec (- n 1) (f accum n))
  ))))
    (body n accum)
)))) ... )
```

One can use function value parameters in recursive functions by simply defining a non-recursive wrapper function.

The author did not know about the Curry paradox and functional recursive combinators (maybe he shouldn't have picked a functional source language) and one can write recursive

expressions like that in λV as well. Below is a (simple) example showing how simple *func* constructs can be (ab-)used to get recursion.

```
(let
  ((sumup (func (self n)
    (if (eq n 0) 0 (+ n (self self (- n 1)))))))
  (sumup sumup 10))
```

But those expressions are not well-typed in our source language. There is no (finite) derivation tree for the well-typedness of the example expression since we define our type rules by substituion, meaning that for recursive function constructs like this one would need an infinitely large derivation tree (independent from whether or not the expression actually terminates). In short: we expect programmers to play nice and use the *rec-func* construct we provide since we can't support arbitrary recursion. Sadly.

There are furtheremore a lot of more uninteresting typing rules for the builtin primitives such as arithmetic or trigonometric functions. We annotate those builtins with types (just some examples in the list below):

- $Vec(I, f) \to Vec(I, f)$, e.g. the unary minus, fract, exp

- $Vec(I, f) \to Num$, e.g. length

- $Vec(b, f) \to Bool$, e.g. any-of

- $Vec(I, f) \times Vec(I, f) \to Vec(I, f)$, e.g. plus

- $Vec(I, f) \times Vec(I, f) \to Num$, e.g. distance

- $Vec(I, T) \times Vec(I, T) \to Vec(I, b)$, e.g. less-than or equal

Note that the generic i must be the same for all parameters/return types. This allows us to just give one generic rule for all of those builtins:

$$\frac{builtin \text{ of type } (\tau_1 \dots \tau_n) \to \tau \qquad \varnothing, \varnothing \vdash e_i : \tau_i}{R, \varnothing \vdash (builtin \; e_1 \dots e_n) : \tau}$$

6   Translation

We model translation as a recursive function *translate*. We use the usual structure (unordered tuples where each value instead gets a name) and member-access notation to keep things somewhat compact.

$$BackEdge ::= (ID, ID*)$$
$$GE ::= (ID, \tau)$$
$$R_i ::= defs : identifier \rightarrow GE, cont : ID$$
$$R_o ::= BackEdge*$$
$$ID ::= \text{SPIR-V numeral ID}$$
$$A ::= (ce*)*$$
$$G ::= identifier \rightarrow GE$$
$$C ::= [ID \mapsto Num]*$$

$$Input ::= \{expr : e, args : A, gen : G, idc : ID, idl : ID, rec : R_i\}$$
$$Output ::= \{code : \vec{I}, consts : C, ido : ID, idc : ID, idl : ID, type : \tau, rec : R_o\}$$

$$translate : Input \rightarrow Output$$

Input.expr is simply the expression to be translated. We extend expressions by an intermediate type generated during translation, an already translation expression holding its ID and its type.

Input.rec and **Output.rec** are information needed to generate recursive functions. Inside a *rec-func* construct, **Input.rec** is the block ID of the continue block (to jump to) and the IDs as well as types (bundled into our new GE type) for the recursive function arguments.

Output.rec is a set of blocks and their respective parameter IDs for recursvie calls (i.e. edges to the continue block).

Input.args models the current argument call stack, pretty much the same way we modeled it for our typing rules. This is once again needed because we can't translate function expressions on their own.

ID represents a SPIR-V ID.

Input.idc is the next usable ID.

Input.idl is the ID of the current code generation block (i.e. label).

Output.ido is the ID holding the result of the translated expression.

Output.idc is the next usable ID after the translation. E.g. if a translation of an expression gets $Input.idc = 42$ as input, uses IDs 42, 43, 44 and 45, it returns 46 as $Output.idc$.

Output.idl is the label the code generation finishes in. This is different iff code generation inserts a new *OpLabel* instruction, i.e. starting a new basic block.

$Output.consts$ is a set of defined constant instructions. In SPIR-V constants cannot be defined inline but have to be defined in a special section before the start of the program, that's why keep them as a separate vector.

$Output.type$ is the type of the generated expression.

Finally, $Output.code$ is the generated SPIR-V instruction vector.

6.1   Translation: utility

We define the function as a set of conditions in which $translate(I) = O$ is defined. The next sections will present a set of conditions for each expression, you can basically imagine each section being one (giant) derivation rule for the defining judgment $translate(I) = O$ with all the conditions as premises.

When generating the SPIR-V module from a full λV program, we define all types in the header (giving them IDs) and can therefore define a function $typeid : \tau \to ID$ that returns the SPIR-V type ID associated with a given type in our formalization. We define $typeid(Rec\tau) = typeid(\tau)$. The expression $typeid(PureRec)$ is intentionally undefined, it is never needed for well-defined programs. Similar to the way we declare types once and can then use them via their IDs during the translation, we also define the constants $true$ and $false$ once and are able to use them during translation. Their translation rules are not shown below since they are therefore trivial (basically a no-op, just returning the ID of the respective constant).

We also need a utility function ioa, modeling the *insert or assign* semantic of a mapping (in our compiler we simply use a hash map):

$$ioa(Mapping, id \mapsto v)(id_c) := \begin{cases} v & \text{for } id = id_c \\ Mapping(id_c) & \text{otherwise} \end{cases}$$

We write $ioa(Mapping, (id \mapsto v)^*)$ as a shortcut for subsequent insertion/replacement,

$$ioa(ioa(\dots ioa(Mapping, id_1 \mapsto v_1)\dots), id_n \mapsto v_n)$$

6.2   Translation: func

This is probably the most interesting (and yet one the most simple) translation rules: to translate a function call we simply translate its body (effectively always inlining the function) and replace all occurrences of function parameter with translations of the expressions bound to them. Of course, this might lead to code bloat, when huge functions or complex expression passes as paremters are inlined. This compiler doesn't care too much and separate optimization passes can still perform common subexpression elimination or similar (they could technically even refactor code that is generated multiple times out into its own function, if possible). The reason we translate functions (or rather: function calls; we can't translate function in itself, remember how they are not even valid expressions since not typed at all) like this even though SPIR-V offers functions is that there is no dynamic dispatch in SPIR-V. Therefore force-inlining basically everything is the only way to get higher-order functions (well, with the restrictions outlined in the beginning).

I.expr $= (\text{func } (id_1 \dots id_n) \text{ e})$
I.args $= (A_r, (e_1 \dots e_n))$
   O $= \text{translate}(\{\text{expr: } e[e_1/id_1]\dots[e_n/id_n], \text{ args: } A_r, \text{ idc: I.idc, idl: I.idl, rec: I.rec}\})$

6.3   Translation: numbers

All constants must be declared in the SPIR-V header, we therefore don't generate code for a constant number but simply return our constant definition (in O.consts) and return the id of the constant.

$$
\begin{aligned}
\text{I.expr} \quad &= num \\
\text{I.args} \quad &= \varnothing \\
\text{O.code} \quad &= \varnothing \\
\text{O.consts} \quad &= [\text{I.idc} \mapsto num] \\
\text{O.idc} \quad &= \text{I.idc} + 1 \\
\text{O.type} \quad &= \text{Num} \\
\text{O.idl} \quad &= \text{I.idl} \\
\text{O.rec} \quad &= \varnothing
\end{aligned}
$$

6.4   Translation: let

$$
\begin{aligned}
\text{I.expr} \quad &= (\text{let } ((id_1 e_1) \ldots (id_n e_n)) \text{ e}) \\
\text{O} \quad &= \text{translate}(\{\text{expr: } e[e_1/id_1] \ldots [e_n/id_n], \text{ args: I.args, idc: I.idc, idl: I.idl, rec: I.rec}\})
\end{aligned}
$$

6.5   Translation: list

This translation rule is only useful (and well-defined) when e_0 isn't an atomic expression such as $func$ or a builtin.

$$
\begin{aligned}
\text{I.expr} \quad &= (e_0 \ e_1 \ldots e_n) \\
\text{nargs} \quad &:= (\text{I.args}, (e_1 \ldots e_n)) \\
\text{O} \quad &= \text{translate}(\{\text{expr: } e_0, \text{ args: nargs, idc: I.idc, idl: I.idl, rec: I.rec}\})
\end{aligned}
$$

6.6   Translation: if

This is the first translation actually generating SPIR-V code. The main complexity in this translation is handling cases where one (or both) of the given branches is just a recursive call (i.e. of type $PureRec$).

$$
\begin{aligned}
\text{I.expr} \quad &= (\text{if } e_c\ e_t\ e_f) \\
O_c \quad &= \text{translate}(\{\text{expr: } e_c,\ \text{args: I.args, idc: I.idc} + 4,\ \text{idl: I.idl, rec: I.rec}\}) \\
O_t \quad &= \text{translate}(\{\text{expr: } e_t,\ \text{args: I.args, idc: } O_c.\text{idc},\ \text{idl: I.idc} + 0,\ \text{rec: I.rec}\}) \\
O_f \quad &= \text{translate}(\{\text{expr: } e_f,\ \text{args: I.args, idc: } O_f.\text{idc},\ \text{idl: I.idc} + 1,\ \text{rec: I.rec}\}) \\
\text{O.consts} \quad &= \text{consts: } O_c.\text{consts } O_t.\text{consts } O_f.\text{consts}
\end{aligned}
$$

$$
\text{O.ido} \quad = \begin{cases} 0 & \text{for } O_f.type = PureRec \wedge O_t.type = PureRec \\ I.idc + 4 & \text{otherwise} \end{cases}
$$

$$
\text{O.idl} \quad = \begin{cases} 0 & \text{for } O_f.type = PureRec \wedge O_t.type = PureRec \\ I.idc + 2 & \text{otherwise} \end{cases}
$$

$$
\text{O.type} \quad = \begin{cases} O_t.type & \text{for } O_f.type = PureRec \\ O_f.type & \text{otherwise} \end{cases}
$$

$$
\begin{aligned}
\text{O.idc} \quad &= O_f.\text{idc} \\
\text{O.rec} \quad &= O_t.\text{rec},\ O_f.\text{rec}
\end{aligned}
$$

O.code is defined as:

	O_c.code	
	OpSelectionMerge (I.idc + 4) None	
	OpBranchConditional O_c.ido (I.idc + 1) (I.idc + 2)	
(I.idc + 0) =	OpLabel	
	O_t.code	
	OpBranch (I.idc + 2)	[**if** O_t.**type** $\neq$ **PureRec**]
(I.idc + 1) =	OpLabel	
	O_f.code	
	OpBranch (I.idc + 2)	[**if** O_f.**type** $\neq$ **PureRec**]
(I.idc + 2) =	OpLabel	[**if** O.**type** $\neq$ **PureRec**]
(I.idc + 3) =	OpPhi typeid(O.type)	[**if** O.**type** $\neq$ **PureRec**]
	O_t.ido (O_t.idl)	[**if** O_t.**type** $\neq$ **PureRec**]
	O_f.ido (O_f.idl)	[**if** O_f.**type** $\neq$ **PureRec**]

6.7   Translation: Computations

The various builtins are intuitively translated to a single instruction. For instance, a translation of the "vec4" builtin, using 4 numbers to construct a *Vec(4,f)* could look like this:

$$
\begin{aligned}
\text{I.expr} \quad &= \text{vec4} \\
\text{I.args} \quad &= ((e_1\ e_2\ e_3\ e_4)) \\
O_0 \quad &= \{\text{idc: I.idc} + 1\} \\
O_i \quad &= \text{translate}(\{\text{expr: } e_i,\ \text{args: } \varnothing,\ \text{idc: } O_{i-1}.\text{idc},\ \text{idl: I.idl, rec: I.rec}\})\ \ldots \\
\text{O.idc} \quad &= O_4.\text{idc} \\
\text{O.ido} \quad &= \text{I.idc} + 1 \\
\text{O.type} \quad &= \text{Vec}(4,\ \text{f}) \\
\text{O.rec} \quad &= \varnothing \\
\text{O.consts} \quad &= O_i.consts\ \ldots \\
\text{O.code} \quad &= (\text{I.idc} + 1) = \text{OpCompositeConstruct typeid(Vec(4, f))}\ O_1.\text{ido }O_2.\text{ido }O_3.\text{ido }O_4.\text{ido}
\end{aligned}
$$

6.8   Translation: rec-func

Translating *rec-func* is by far the most complicated translation. Since most SPIR-V runtimes don't allow recursion (GPUs traditionally don't have a stack) we have to unroll it into a loop. That is the reason we only support this limited form of recursion in our source language. We provide the frame, including all basic blocks and SSA phi functions for recursive calls from within the function body, i.e. (after translation) jumps back to the begin of the loop body (via a separate continue block SPIR-V needs) from within the loop body.

$$
\begin{aligned}
\text{I.expr} \;&= (\text{rec-func } (id_1 \ldots id_n) \text{ e}) \\
\text{I.args} \;&= ((e_1 \ldots e_n))
\end{aligned}
$$

$$
\begin{aligned}
\text{hb} \;&:= \text{I.idc} + 0 \\
\text{lb} \;&:= \text{I.idc} + 1 \\
\text{cb} \;&:= \text{I.idc} + 2 \\
\text{mb} \;&:= \text{I.idc} + 3 \\
O_0 \;&:= \{\text{idc: I.idc} + 4\} \\
O_i \ldots \;&:= \text{translate}(\{\text{expr: } e_i, \text{ args: } \varnothing, \text{ idc: } O_{i-1}.\text{idc, idl: I.idl, rec: I.rec}\}) \ldots \\
D_b \;&:= \text{ioa}(\text{I.rec.defs}, (id_i \mapsto (O_n.idc + i, O_i.\text{type})) \ldots) \\
O_b \;&:= \text{translate}(\{\text{expr: e, idc: } O_n.idc + 2n, \text{ idl: lb, rec: cont: cb, defs: } D_b\})
\end{aligned}
$$

$$
\begin{aligned}
O_b.\text{type} \;&= \text{Rec } \tau \\
O_b.\text{rec} \;&= ((block_1, bid_1^1 \ldots bid_n^1) \ldots (block_m, bid_1^m \ldots bid_n^m))
\end{aligned}
$$

$$
\begin{aligned}
\text{O.rec} \;&= \varnothing \\
\text{O.consts} \;&= O_b.\text{consts } O_i.\text{consts} \ldots \\
\text{O.ido} \;&= O_b.\text{ido} \\
\text{O.idl} \;&= \text{mb} \\
\text{O.type} \;&= \tau
\end{aligned}
$$

Furthermore, *O.code* is defined as:

	OpBranch hb	
	O_i.code	[for i = 1..n]
hb =	OpLabel	
$(O_n$.idc + i$)$ =	OpPhi typeid(O_i.type) O_i.ido *I*.idl $(O_n$.idc + n + i$)$ cb	[for i = 1..n]
	OpLoopMerge mb cb None	
	OpBranch lb	
lb =	OpLabel	
	O_b.code	
	OpBranch mb	
cb =	OpLabel	
$(O_n$.idc + n + i$)$ =	OpPhi typeid(O_n)	[for i = 1..n]
	$block_k \; bid_i^k$	[for k = 1..m]
	OpBranch hb	
mb =	OpLabel	

6.9   Translation: rec

The definition of this *rec* translation only makes sense together with the translation of *rec-func*. Interesting here is that this translation does not return a value or block ID (dummy value 0) in O.ido, O.idl, since neither is defined. For well-typed expressions this will never again be needed during the translation since the result of this expression can't be used anyways, that is the idea of tail recursion. With our current type system, every *rec* expression must be wrapped immediately into an *if* expression.

$$
\begin{aligned}
\text{I.expr} &= \text{rec} \\
\text{I.args} &= ((e_1 \ldots e_n)) \\
O_0 &= \{\text{idc: I.idc}\} \\
O_i &= \text{translate}(\{\text{expr: } e_i, \text{ args: } \varnothing, \text{ idc: } O_{i-1}.\text{idc, idl: I.idl, rec: } \varnothing\}) \ldots
\end{aligned}
$$

$$
\begin{aligned}
\text{O.idc} &= O_n.\text{idc} \\
\text{O.ido} &= 0 \\
\text{O.idl} &= 0 \\
\text{O.type} &= \text{PureRec} \\
\text{O.consts} &= O_i.consts \ldots \\
\text{O.rec} &= (\text{I.idl}, (O_i.\text{ido} \ldots)) \\
\text{O.code} &= O_i.code \ldots \quad \text{OpBranch I.rec.cont}
\end{aligned}
$$

6.10   Translation: identifier

We replace all identifiers except the function parameters in *rec-func* constructs, since those cannot be replaced by source-lanauge expressions but must simply use the results from the *OpPhi* instructions in the beginning of the *rec-func* loop body.

$$
\begin{aligned}
\text{I.expr} &= identifier \\
\text{I.args} &= \varnothing \\
\text{I.rec}(identifier) &= (ido, \, tau) \\
\text{O} &= \{\text{code: } \varnothing, \text{ consts: } \varnothing, \text{ ido: } ido, \text{ idc: I.idc, idl: I.idl, type: } \tau, \text{ rec: } \varnothing\}
\end{aligned}
$$

6.11   Module creation

A valid SPIR-V module requires more than just the instruction vector. We have to add a header, define an entry point, the used extensions, all types and constants (generated/used by the translation) as well as the one output variable (representing the framebuffer output, i.e. where we will write the resulting value in the end) which must be annotated with an *OpDecorate* instruction that connects it to the framebuffer output. Then we define the entry point function and the first block using *OpLabel*. After that we insert the code generated by the translation of the source expression. In the end we have to store the SPIR-V value in the global output using *OpStore* on O.ido from our translation. Then we end our main function using *OpReturn* and *OpFunctionEnd*.

We statically use IDs for types and the entry block. We therefore define the first usable ID (i.e. the number of IDs we statically use in the header + 1) as *startid* and the ID of the

first basic block (*OpLabel*) in our main function as *startblock*. Both of these are required in the translation of our source expression.

7 Correctness

Note: the proof regarding the tail recursion (i.e. involving rec) is not completely formalized and has some problems. I had trouble finding a formalization in which this can obviously and easily be proven.

We want to prove whole program correctness. Valid programs are expressions that are well-typed with no recursive context and no arguments and have a type τ_o that qualifies as observation, i.e. $Vec(i, f)$, since only those values can be returned (written into a framebuffer) by a fragment shader. We furthermore define the utility function $init_S : (C \times \vec{I}) \to M$ that returns a SPIR-V memory object initialized with the constants from C and the block mappings from the full-program vector given in the $\vec{I}$ argument. Another utility function $eval_S : (M \times ID) \to Vec(1, f)$ models the full evaluation of the program loaded into the present SPIR-V memory and returns the observation mapped to the given ID after program execution according to the SPIR-V operational semantics we outlined. It basically starts execution at the block *startblock* and applies the rules from the SPIR-V operational semantics until no instructions are left and then returns the value present in the memory for the given ID.

Another utility function to connect definitions and substitution: $subst : (e \times D) \to e$ is defined as $e[e_i\ id_i]\dots$ for every $id_i \mapsto e_i$ mapping in the given definitions mapping.

The correctness theorem looks like this:

When we have a well-typed program e with
$\varnothing, \varnothing \vdash e : \tau_o$ and
$eval_\lambda(e) = o$ in λV and
$translate(\{expr : e, args : \varnothing, idc : startid, idl : startblock, rec : \varnothing\}) = O$,
then $eval_S(init_S(O.consts, O.code), O.ido) = o$.

To actually prove this we need to strengthen the induction argument, accounting for arguments and types:

When we have a well-typed program e with
$\varnothing, ((e_1^1 \dots) \dots (e_n^1 \dots)) \vdash e : \tau$ and
$eval_\lambda(((e\ e_n^1 \dots) \dots)e_1^1 \dots) = o$ in λV and
$translate(\{expr : e, args : ((e_1^1 \dots) \dots (e_n^1 \dots)), idc : startid, idl : startblock, rec : \varnothing\}) = O$,
then
$O.type = \tau$,
if $O.ido \neq 0$, it has type $idtype(\tau)$;
and $eval_S(init_S(O.consts, O.code), O.ido) = o$.

We did not formally define the type system for SPIR-V but it should be obvious from its text specification.

We prove this theorem by induction over the well-typedness, i.e. the type derivation tree of expression e. We can start with the simple cases, even without additional Lemmas.

7.1   Correctness: numbers

When e is a number num it obviously evaluates to itself as observation. Looking up its translation rule, we see that it translates to no code at all. But since $eval_S(init_S(O.consts, O.code), O.ido)$ returns the value of $O.ido$ — in case of the translation of a constant that's simply the ID we gave the constant — in the memory after execution — which in our case is the same memory as before the execution — we simply get num since that's the constant we added in our translation, that gets loaded into initial memory by $init_S$. We could trivially formalize this proof given formalizations of $eval_S$ and $init_S$ but the proof would just trivially look like described here.

7.2   Correctness: computations

Since we defined all builtins to just have the semantics of their SPIR-V counterparts (see the example for the operational semantics of OpFAdd, all we do is basically saying "this function behaves like specified in SPIR-V"), the correctness proof for those is trivial. We use the induction hypothesis for the arguments passed to the builtins.

To formally prove this we would have to strengthen our induction argument, stating that $O.type$ is the same as the type of e and that the $O.ido$ is a value of this type as well. Otherwise we can't guarantee that the types actually match and that the instruction is valid. But this can be verified separately or via a Lemma by just looking at each translation rule.

7.3   Correctness: list

Premises:

- $e = (e_0\, e_1 \ldots e_n)$, well typed under arguments A_r

- Per induction hypothesis (see typing rule $t\text{-}app$), we know that our strengthened theorem holds for e_0 with arguments $(A_r, (e_1 \ldots e_n))$.

This translation is basically just a utility step, moving arguments from the list into our translation argument stack (i.e. we translate e_0 with new argument stack $(A_r, (e_1 \ldots e_n))$). The proof for the correctness of the translation is immediately given transitively by our induction hypothesis, since $(e_0 e_1 \ldots e_n)$ with arguments A_r has the same type as e_0 with arguments $(A_r, (e_1 \ldots e_n))$ and furthermore e with applied argument stack A_r (as seen in the premise of our strengthened theorem) is syntactically exactly the same as e_0 with unwrapped argument stack $(A_r, (e_1 \ldots e_n))$, meaning they both obviously evaluate to the same observation.

7.4   Correctness: func

Premises:

- $e = (func\ (id_1 \ldots id_n)\ e_b)$, well typed under arguments $(A_r, (e_1 \ldots e_n))$

- Per induction hypothesis (see typing rule *t-app*), we know that our strengthened theorem holds for $e_b[e_1/id_1] \ldots [e_n/id_n]$ with arguments (A_r).

Our translation is defined as $O = translate(\{expr:\ e[e_1/id_1] \ldots [e_n/id_n],\ args:\ A_r,\ idc:$ *I.idc, idl: I.idl, rec: I.rec*$\})$. Since per typing rules, e and $e[e_1/id_1] \ldots [e_n/id_n]$ have the same type and per operational semantics, e can be reduced to $e[e_1/id_1] \ldots [e_n/id_n]$ i.e. both evaluate to the same value, we can just directly use our induction hypothesis to prove that this translation step is correct.

7.5   Correctness: let

Premises:

- $e = (let\ (\ldots (id_i\ e_i) \ldots)\ e_b)$, well typed under arguments A

- Per induction hypothesis our strengthened theoerem holds for $e_b[e_i/id_i] \ldots$, with the same arguments A

Our translation once again allows us to directly use the induction hypothesis to prove correctness, since we simply translate it to the translation of $e_b[e_i/id_i] \ldots$ with the same argument. The operational semantics give us that e can be reduced to $e_b[e_i/id_i] \ldots$ i.e. both evaluate to the same value. The typeing rules guarantee us that both have the same type.

7.6   Correctness: if

Our premises:

- $e = (if\ e_c\ e_t\ e_f)$, well typed under arguments $((e_1^1 \ldots) \ldots (e_n^1 \ldots))$.

- Per induction hypothesis we can assume the correctness theorem for e_c, e_t and e_f

- There are two possible cases for the reduction of e: if e_c evaluates to true, e evaluates to whatever e_t evaluates to otherwise whatever e_f evaluates to (given the arguments)

We step through the code generated by the translation of e using our SPIR-V operational semantics:
First, O_c.code is executed, per induction hypothesis this loads the observation of $eval_\lambda(e_c)$ into the memory at id O_c.ido. Based on our typing rules and the (strengthened) induction hypothesis applied to e_c we know that O_c.ido must be of type bool and the following *OpBranchCondition* can therefore be reduced. *OpSelectionMerge* has no effect, and can effectively be ignored in our reduction, as previously described. If O_c.ido was true, we will execute O_t.code, otherwise O_f.code. Per induction hypothesis, both are guaranteed to load

the observation $eval_\lambda(e_t)$ (or $eval_\lambda(e_f)$, respectively) into O_t.ido (or O_f.ido, respectively), assuming that the types of the branch expressions are not *PureRec* (we get to that case below). After that, control always branches to the block *I.idc + 2*. All that is done there is to select the computed value via an *OpPhi* instruction. This value is the return value of the translation.

If one of the branch expressions is of type PureRec, that means per typing rules that it is (potentially nested in some called *func* or *let* expressions) an expression *(rec ...)* and can therefore not be reduced (we have no reduction rule for *rec*). So if this branch is executed, we don't have to fulfill anything for our output (the premises of the correctness theorem are not fulfilled). On the other hand, this means (per induction hypothesis) that its translation will return O.type *PureRec* and can return no output id, so we must not use its *O.ido* or *typeid(O.type)*. It furthermore will not follow linear flow, i.e. end the current code block, so the next instruction afterwards must be *OpLabel*, starting a new block, for our SPIR-V code vector to be well defined.

The returned value *O.ido* is either 0 (in the case when both branches are of type *PureRec*) or of type *typeid(O.type)*. When either of the branches is not of type *PureRec*, a new ID will be returned as output, defined by the *OpPhi* instruction in the last block. This instruction is of type *typeid(O.type)*. Furthermore, it is well-defined since all passed parameters are of the same type (remember that $typeid(Rec\tau) = typeid(\tau)$), that is guaranteed by our typing rules and the induction hypothesis applied to the (non-*PureRec*) branch expressions.

7.7   Correctness: rec-func and rec

This is by far the most complicated proof. We have to prove *rec-func* and *rec* together because they depend on each other and because *rec* in itself cannot be reduced, it will always be substituted instead by the *rec-func* expression it is contained in.

Our premises:

- $e = (\text{rec-func } (id_1 \ldots) \ e_b)$

- The arguments for the well-typedness of e are $((e_1 \ldots e_n))$.

- e reduces to $e_b[e_1/id_1] \ldots [e_n/id_n][(\text{rec-func } (id_1 \ldots id_n) \ e)/rec]$

- Per induction hypothesis we know that our correctness theorem holds for $e_b[e_1/id_1] \ldots [e_n/id_n]$, translated with no arguments.

- Per induction hypothesis we also know that our correctness theorem holds for all the arguments e_i.

We reduce the code generated by the translation of e according to the SPIR-V semantics we defined (writing a derivation tree for this would require a lot of space):
Since we know the correctness theorem holds for the function arguments, we know that $O_i.code$ in the *I.idl* block indeed produces the same observations produced by $eval_\lambda(e_i)$. There is a reduction rule giving us in addition basically that $(\ldots e_i \ldots)$ is the same as $(\ldots o \ldots)$, given that $eval_\lambda(e_i) = o$. When we get to *OpPhi* in block *hb*, O_i.ido will be

selected since we jumped to the block from *I.idl*. We ignore the OpLoopMerge instruction and jump to *lb*, where O_b gets executed. The only way O_b won't reach the *OpBranch* ending the block is if execution gets to the translation of *rec* (no other translations will ever branch out of linear flow) associated with the current *rec-func* frame. Per opertional semantics, *rec* in that case will get substituted by the original *rec-func* expression.

Our translation of *rec* will generate the code for all its arguments (here we can use the induction hypothesis for *rec* to argue that the generated translations produce the same observations) and branch to *cb*. In *cb*, those arguments will be selected via *OpPhi*, and control will branch to *hb* again, where those same arguments (transitively through the previous *OpPhi*) function will be selected for the arguments used by O_b.*code*, where control branches subsequently.

This means, given that our correctness theorem holds for $e_b[e_1/id_1]\ldots[e_n/id_n]$ (which we know per induction hypothesis for *rec-func*), we know that once control reaches the *OpBranch mb* instruction, the observation stored in O_b.ido (which is what is returned in O.ido) should indeed hold the observation e can be reduced to.

The well-typedness of our translation (as required by our strenghtened induction argument) is similarly guaranteed by the our induction hypothesis, i.e. the well-typedness of the function body and its arguments.

This argumentation isn't yet enough for a formal proof but with enough time (and maybe some needed re-formalization) it should be possible to transform this reasoning into a formal proof, given that most of our explanation is just applying the SPIR-V operational semantics we defined.

8   Conclusion

We presented a translation for a language with functional elements such as (tail) recursion and higher-order functions (with restrictions) to SPIR-V, a language without dynamic memory allocation, recursion, a stack frame or dynamic dispatch.

Initially, we modeled translation without using our "magical" context-sensitive substitution. Instead, we used an additional translation input parameter holding all definitions and passing the environment for each argument tuple. While that was the last step to basically make translation just a one-to-one formal model of our code, it made the formalization much more complicated, requiring multiple helper functions, an additional strenghtening of our correctness induction hypothesis and in the end everything just boiled down again to a Lemma saying that translation of an expression with mappings given in the "definitions" argument is the same as just substituting them.

We still believe though that it would be worth to verify this Lemma (i.e. verify that substitution is indeed done correctly by our code) since this task was not trivial while writing the compiler.

Further work includes extending this formal model from the subset to a full-blown shader language with inputs, multiple outputs as well as buffer and image read/write (modeled as inputs or observations, respectively). We would also like to explore how far other functional features such as sum types and pattern matching can be translated to SPIR-V. To actually make this compiler presented here usable, we would have to explore in how far the existing

optimization tools for SPIR-V binaries are able to minimize the code bloat and performance problems introduced by our translation. The *spirv-opt* program delivered promising results in first experiments. It would be interesting to see if there are optimizations it cannot perform that could instead be done in our compiler. On the other hand, using additional (unverified) optimization passes obviously might mess with the outlined correctness of our compiler.

Furthermore, our SPIR-V binaries (like the output of all shader languages known to me) can contain undefined behavior. I want to look into a general SPIR-V compiler pass that just transforms binaries into a *debug* mode, inserting runtime checks for undefined behavior conditions, writing detected problems to a storage buffer object (e.g. writing an ID associated with the condition and the source location). There already exists a framework for SPIR-V compiler/optimization passes (that even implements outputting debug values to a storage buffer) in the SPIRV-Tools project repository [7] (of which *spirv-opt* is a part which could probably be easily extended by this kind of pass.

---

[7]https://github.com/KhronosGroup/SPIRV-Tools